

MD Sohrab Ahmed

+60 1112625189 | sohrabahmed@outlook.com | [linkedin.com/in/sohrabahmed](https://www.linkedin.com/in/sohrabahmed) | Kuala Lumpur, Malaysia

SUMMARY

Lead Data Engineer with 11+ years of experience designing and building scalable data platforms, real-time data pipelines, and modern lakehouse architectures across AWS, Azure, Cloudera, and Databricks.

Expert in Databricks Lakehouse, Apache Spark, Kafka, and Flink with strong focus on data quality, governance, cost optimization, and real-time analytics. Proven experience in leading enterprise data warehouse migrations, building metadata-driven frameworks, and enabling large-scale data platform adoption across engineering teams.

EXPERIENCE

Senior Data Engineer | SEEK, Kuala Lumpur, Malaysia | December 2024 – Present

Skills: Databricks (DQX, Unity Catalog, Genie), AWS (MSK, S3, Redshift), Apache Flink, Apache Kafka, Spark Structured Streaming, Python, Claude Code, Braze, Cost Optimization, Mentoring

- Architected a high-throughput **real-time data platform** using Apache Flink, Kafka (AWS MSK), and Databricks, reducing data latency from hours to sub-minute for critical business decision-making
- Architected AI-driven workflows via **Claude Code** to automate end-to-end ETL lifecycles, DQM checks, and job deployments. Implemented closed-loop validation triggering Git updates and automated PRs, significantly enhancing engineering velocity and data integrity
- Engineered a **centralized data quality and observability framework** using Databricks DQX, implementing automated rule orchestration and anomaly detection across all global pipelines
- Enabled self-service analytics via **Databricks Genie**, utilizing AI-driven semantic layers to provide non-technical stakeholders with natural language query capabilities
- Defined platform standards and reusable frameworks, successfully onboarding 5+ engineering squads to the Databricks Lakehouse and accelerating feature delivery cycles by ~20%
- Spearheaded **FinOps** initiatives by designing S3 lifecycle policies (Glacier Instant Retrieval) and cluster overrides, significantly optimizing storage and compute spend
- Collaborate as a technical lead with product, marketing and analytics teams to modernize legacy architecture and ensure a scalable, governed data ecosystem

Senior Data Engineer | FedEx, Hyderabad, India | August 2023 – December 2024

Skills: Microsoft Azure, Data Factory, Databricks, Synapse Analytics, SQL, PySpark, Scala, Delta Lake, DevOps, CI/CD, SAS

- Built metadata-driven ABC (Audit, Balance, Control) framework improving data quality, reliability, and operational efficiency
- Led re-architecture of MARS platform using Azure Databricks and Unity Catalog, improving scalability, governance, and security
- Migrated SAS Enterprise APAC reporting workloads to Synapse serverless, achieving ~30% cost reduction for 500+ users
- Implemented data lineage and dependency tracking across systems ensuring transparency and compliance across data lifecycle
- Established governance policies and CI/CD standards to improve code quality and accelerate delivery cycles
- Worked in SAFe Agile environment aligning engineering execution with business priorities

Senior Data Engineer | PPG Industries, Subang Jaya, Malaysia | November 2019 – December 2022

Skills: Azure, Databricks, Delta Lake, Snowflake, Power BI, Synapse, ADF, SQL Server, DevOps

- Built scalable Delta Lake platform supporting ACID transactions, schema evolution, time travel, and SCD Type 1/2 processing
- Developed metadata-driven ingestion and integration pipelines across Azure and Snowflake ecosystems
- Reduced data defects by 84% through proactive data quality rules and monitoring frameworks
- Designed data consumption layer using Snowflake and Power BI for 200GB+ datasets
- Implemented CI/CD pipelines for Databricks and ADF, reducing manual effort by 90%
- Optimized Power BI and Analysis Services refresh processes improving performance without additional cost

Big Data Engineer | Neustar, Hyderabad, India | December 2018 – October 2019

Skills: Spark, Hive, Hadoop, Cloudera, Airflow, Scala, S3

- Built automated data quality framework to detect anomalies and ensure data consistency
- Developed centralized data catalog with standard naming conventions improving governance
- Designed dashboards for pipeline health, refresh rate, and data availability KPIs

Data Engineer | Sphota, Bengaluru, India | April 2018 – December 2018

Skills: Hadoop, Hive, Spark, Teradata, Azkaban, Power BI

- Developed automated data ingestion framework with scheduling, monitoring, and alerting
- Built low-latency query solution for clickstream analytics
- Implemented GDPR-compliant transformations using custom UDFs

Programmer Analyst | Cognizant, Gurugram, India | June 2015 – March 2018

Client: American Express

Skills: Hadoop, Hive, Sqoop, Spark, Scala, SQL

- Developed ETL pipelines using Hive and Sqoop for enterprise data ingestion
- Built custom UDFs and UDAFs in Scala for business logic implementation
- Supported large-scale risk analytics data processing on Hadoop ecosystem

CERTIFICATIONS

- Databricks Certified Data Engineer Professional
- Databricks Certified Platform Architect
- Microsoft Certified: Azure Data Engineer, Azure Developer, Power BI Data Analyst, Azure Fundamentals
- Cloudera Certified Associate (CCA) Spark and Hadoop Developer
- SAFe 6 Practitioner
- Microsoft Certified Trainer (MCT)

EDUCATION

- M.Tech – Computing Systems and Infrastructure
BITS Pilani, India | 2023
- B.Tech – Computer Science and Engineering
Uttar Pradesh Technical University, India | 2014